

Production Monitoring Checklist

Web3Geeks Ticket-Triage Agent

What to track

Metric	What it tracks	Alert threshold	Re-evaluation cadence
Error rate	% of /triage or /resume calls returning 5xx or agent execution failed	> 2% over 15 min	Weekly, plus after any prompt/rule change
Latency (p50/p95)	End-to-end run time per ticket, logged per request	p95 > 2s (once real LLM calls replace mocks)	Weekly
Cost drift	Token usage estimate per run (real cost once live LLM calls are wired in)	> 20% week-over-week increase with no volume increase	Weekly
Output quality drift	Sample and manually re-score N tickets/week against the eval rubric (task success, accuracy, tone, safety)	Any criterion average drops > 1 point vs baseline	Weekly spot-check; full re-eval monthly
No-KB-match rate	% of tickets where retrieve_kb returns matched=false (proxy for KB coverage gaps)	> 15% of weekly volume	Weekly
Human-approval turnaround	Time between awaiting_human_approval and resume	> 4 business hours median	Weekly
Safety gate integrity	% of refund/cancel/billing/P1 tickets that actually passed through human_checkpoint	Must be 100% -- any miss is a P1 bug	Every deploy + weekly audit

Logging already in place

- Every /triage and /resume call logs: ticket_id, input preview, category, priority, needs_approval, status, latency_ms, a token-usage estimate, and any error -- written as one JSON line per event to logs/agent.log.
- Each agent run also carries its own internal event log (per-node timestamps) inside the returned state, useful for step-by-step debugging of a single ticket.

Escalation path

Error-rate or safety-gate alerts page the on-call engineer immediately. Latency, cost-drift, and quality-drift alerts go to a weekly review channel rather than paging -- these are trend signals, not outages.